

MODEL EVALUATION

1. Logistic Regression

A linear model used for binary classification, making predictions based on the weighted sum of input features. It performs well when the relationship between features and target is linear but struggles with complex data.

- Accuracy: 0.498
- Precision/Recall/F1-score
 - ✓ Positive class (1): High recall (0.96) but poor precision (0.50)
 - ✓ Negative class (0): Very low precision (0.45) and recall (0.03)
- Performance Summary
 - ✓ Model leans heavily towards predicting the positive class.
 - ✓ Poor at handling class imbalance or capturing the negative sentiment accurately.

2. Naive Bayes

A probabilistic classifier based on Bayes' theorem, assuming feature independence. It is fast and effective for text classification but can be limited by this strong independence assumption.

- Accuracy: 0.510
- Precision/Recall/F1-score
 - ✓ Positive class (1): Decent precision (0.51) and recall (0.76).
 - ✓ Negative class (0): Poor recall (0.26).
- Performance Summary
 - ✓ Performs marginally better than Logistic Regression.
 - ✓ It's better at recognizing positive sentiment but still struggles with negative class detection.

3. Support Vector Machine (SVM)

A robust model that finds the best hyperplane to separate data points, suitable for high-dimensional data but slow with large datasets.

- Accuracy: 0.498
- Precision/Recall/F1-score
 - ✓ Positive class (1): High recall (0.96) but poor precision (0.50).
 - ✓ Negative class (0): Precision (0.46), recall (0.03) are very low.
- Performance Summary
 - ✓ Similar to Logistic Regression, biased towards predicting the positive class.
 - ✓ Struggles with negative sentiment.

4. Random Forest

An ensemble model that builds multiple decision trees and outputs the majority class. It handles non-linear relationships well but may be overfit without proper tuning.

- Accuracy: 0.616
- Precision/Recall/F1-score
 - ✓ Better balance across both classes (e.g., 0.60 precision for negative class, 0.63 for positive).

- ✓ Still struggles with class imbalance, with some decrease in recall for the positive class.
- Performance Summary
 - ✓ Performs better than Naive Bayes, Logistic Regression, and SVM.
 - ✓ Still not ideal but captures both positive and negative sentiments reasonably well.

5. LSTM (Long Short-Term Memory)

A deep learning model designed to learn dependencies in sequential data, particularly effective for text-based tasks like sentiment analysis, though it requires more training time and resources.

- Accuracy: 0.775 (Validation accuracy: 0.788).
- **Loss:** 0.465 (Training) / 0.447 (Validation).
- Performance Summary
 - ✓ Best performing model with an accuracy of 77.5%.
 - ✓ The validation accuracy of 78.8% indicates strong generalization to unseen data.
 - ✓ LSTM, being a deep learning model, excels at capturing long-term dependencies in text, making it well-suited for sentiment analysis.

Overall Comparison

Model	Accuracy	Strengths	Weaknesses
Logistic Regression	0.498	Simple to implement	Poor handling of negative sentiment
Naive Bayes	0.51	Good for text data	Struggles with negative class
SVM	0.498	Works well for binary classification	Biased towards positive class
Random Forest	0.616	Captures both classes better	Still suffers from some imbalance
LSTM	0.775	Captures complex text dependencies	Requires more computational resources

Conclusion: LSTM is the Best Model

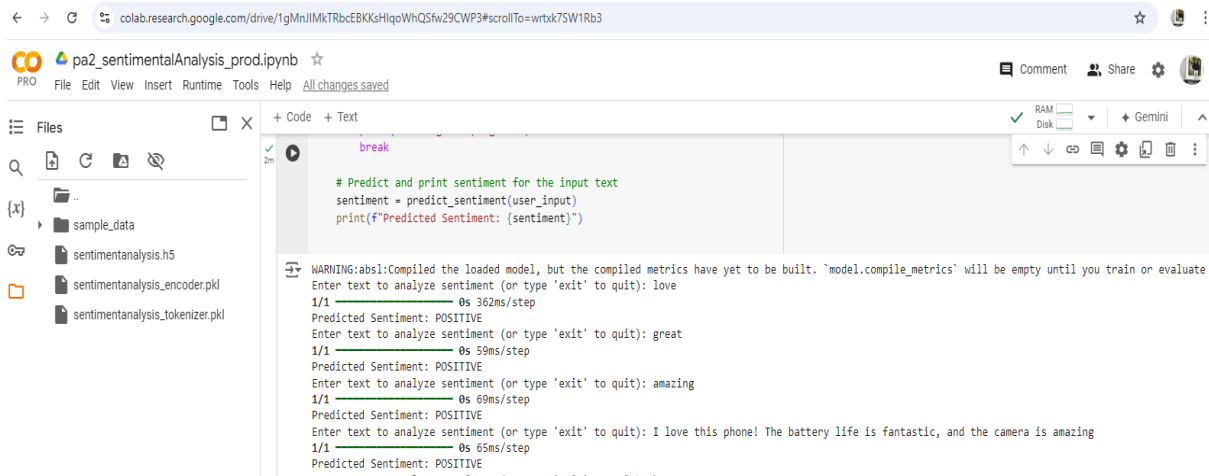
LSTM outperforms all the other models, achieving a significantly higher accuracy of 77.5%. Its ability to model sequential data and understand long-term dependencies makes it ideal for sentiment analysis. While traditional ML models like Logistic Regression and SVM struggle with class imbalance and context, LSTM handles these challenges effectively. Hence, Prediction and Inference Module is based on LSTM model.

TESTING

Note: Recommend running test on google Colab.

Sample output screenshots.

a. Positive Tweets Scenario



The screenshot shows a Google Colab notebook titled 'pa2_sentimentalAnalysis_prod.ipynb'. The code cell contains a function to predict sentiment for input text. The output shows a series of prompts and predictions for positive tweets.

```
# Predict and print sentiment for the input text
sentiment = predict_sentiment(user_input)
print(f"Predicted Sentiment: {sentiment}")
```

WARNING:absl:Compiled the loaded model, but the compiled metrics have yet to be built. 'model.compile_metrics' will be empty until you train or evaluate

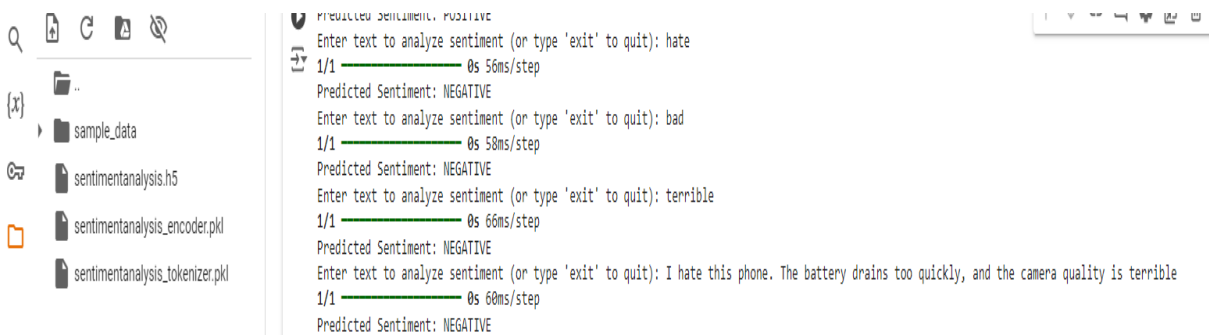
Enter text to analyze sentiment (or type 'exit' to quit): love
1/1 ————— 0s 362ms/step
Predicted Sentiment: POSITIVE

Enter text to analyze sentiment (or type 'exit' to quit): great
1/1 ————— 0s 59ms/step
Predicted Sentiment: POSITIVE

Enter text to analyze sentiment (or type 'exit' to quit): amazing
1/1 ————— 0s 69ms/step
Predicted Sentiment: POSITIVE

Enter text to analyze sentiment (or type 'exit' to quit): I love this phone! The battery life is fantastic, and the camera is amazing
1/1 ————— 0s 65ms/step
Predicted Sentiment: POSITIVE

b. Negative Tweets Scenario



The screenshot shows the same Google Colab notebook as in (a), but with negative tweets. The output shows a series of prompts and predictions for negative tweets.

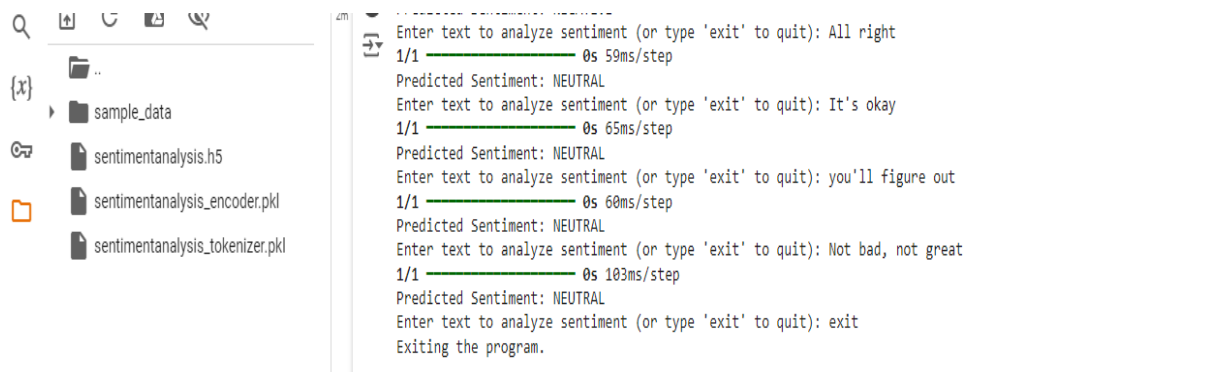
Enter text to analyze sentiment (or type 'exit' to quit): hate
1/1 ————— 0s 56ms/step
Predicted Sentiment: NEGATIVE

Enter text to analyze sentiment (or type 'exit' to quit): bad
1/1 ————— 0s 58ms/step
Predicted Sentiment: NEGATIVE

Enter text to analyze sentiment (or type 'exit' to quit): terrible
1/1 ————— 0s 66ms/step
Predicted Sentiment: NEGATIVE

Enter text to analyze sentiment (or type 'exit' to quit): I hate this phone. The battery drains too quickly, and the camera quality is terrible
1/1 ————— 0s 60ms/step
Predicted Sentiment: NEGATIVE

c. Neutral Tweets Scenario



The screenshot shows the same Google Colab notebook as in (a) and (b), but with neutral tweets. The output shows a series of prompts and predictions for neutral tweets.

Enter text to analyze sentiment (or type 'exit' to quit): All right
1/1 ————— 0s 59ms/step
Predicted Sentiment: NEUTRAL

Enter text to analyze sentiment (or type 'exit' to quit): It's okay
1/1 ————— 0s 65ms/step
Predicted Sentiment: NEUTRAL

Enter text to analyze sentiment (or type 'exit' to quit): you'll figure out
1/1 ————— 0s 60ms/step
Predicted Sentiment: NEUTRAL

Enter text to analyze sentiment (or type 'exit' to quit): Not bad, not great
1/1 ————— 0s 103ms/step
Predicted Sentiment: NEUTRAL

Enter text to analyze sentiment (or type 'exit' to quit): exit
Exiting the program.

MODEL EVALUATION RESULTS

Logistic Regression

```
n_iter_1 = _check_optimize_result(
Logistic Regression Accuracy: 0.497759375
      precision    recall  f1-score   support

      0       0.45      0.03      0.06    159494
      1       0.50      0.96      0.66    160506

 accuracy          0.50    320000
  macro avg       0.47    0.50    0.36    320000
 weighted avg     0.47    0.50    0.36    320000
```

Naive Bayes

```
y = column_or_1d(y, warn=True)
Naive Bayes Accuracy: 0.510165625
      precision    recall  f1-score   support

      0       0.52      0.26      0.34    159494
      1       0.51      0.76      0.61    160506

 accuracy          0.51    320000
  macro avg       0.51    0.51    0.48    320000
 weighted avg     0.51    0.51    0.48    320000
```

Support Vector Machines (SVM)

```
y = column_or_1d(y, warn=True)
SVM Accuracy: 0.4989625
      precision    recall  f1-score   support

      0       0.46      0.03      0.06    159494
      1       0.50      0.96      0.66    160506

 accuracy          0.50    320000
  macro avg       0.48    0.50    0.36    320000
 weighted avg     0.48    0.50    0.36    320000
```

Random Forest

```

return fit_method(estimator, *args, **kwargs)
Random Forest Accuracy: 0.61578125

```

	precision	recall	f1-score	support
0	0.60	0.68	0.64	159494
1	0.63	0.55	0.59	160506
accuracy			0.62	320000
macro avg	0.62	0.62	0.61	320000
weighted avg	0.62	0.62	0.61	320000

Deep Learning models like LSTM (Long Short-Term Memory)

```

Epoch 1/8
1125/1125 — 293s 257ms/step - accuracy: 0.7300 - loss: 0.5298 - val_accuracy: 0.7784 - val_loss: 0.4652 - learning_rate: 0.0010
Epoch 2/8
1125/1125 — 289s 257ms/step - accuracy: 0.7617 - loss: 0.4874 - val_accuracy: 0.7833 - val_loss: 0.4564 - learning_rate: 0.0010
Epoch 3/8
1125/1125 — 289s 257ms/step - accuracy: 0.7674 - loss: 0.4785 - val_accuracy: 0.7860 - val_loss: 0.4530 - learning_rate: 0.0010
Epoch 4/8
1125/1125 — 289s 257ms/step - accuracy: 0.7710 - loss: 0.4735 - val_accuracy: 0.7869 - val_loss: 0.4503 - learning_rate: 0.0010
Epoch 5/8
1125/1125 — 289s 257ms/step - accuracy: 0.7732 - loss: 0.4695 - val_accuracy: 0.7876 - val_loss: 0.4483 - learning_rate: 0.0010
Epoch 6/8
1125/1125 — 290s 258ms/step - accuracy: 0.7730 - loss: 0.4691 - val_accuracy: 0.7880 - val_loss: 0.4486 - learning_rate: 0.0010
Epoch 7/8
1125/1125 — 290s 257ms/step - accuracy: 0.7751 - loss: 0.4669 - val_accuracy: 0.7893 - val_loss: 0.4463 - learning_rate: 0.0010
Epoch 8/8
1125/1125 — 289s 257ms/step - accuracy: 0.7751 - loss: 0.4654 - val_accuracy: 0.7884 - val_loss: 0.4478 - learning_rate: 0.0010

```

Plot Accuracy and Loss Curves (LSTM)

